HIMANSHU PARIHAR

GenAI Engineer

himanshuparihar853@gmail.com | +91 9319264679 | linkedin.com/in/himanshuuk | github.com/Himanshu-18-P

GenAI / LLM Engineer with 2+ years of experience building **production-grade RAG and agentic AI systems at scale**. Skilled in designing **low-latency, cost-efficient LLM applications** using Groq, LangChain, and vector databases. Experienced in improving response reliability through **evaluation frameworks and feedback loops**, with a strong focus on **scalable, monitored, and high-performance AI systems**.

EXPERIENCE

HiVoco — GenAI Engineer , New Delhi, India | Nov 2023 – Present

• Built multiple **production-grade GenAI systems** serving **600K+ users**, processing **200+ concurrent AI workloads**, and supporting high-scale real-time applications

• Developed an AI-powered **video generation platform** for marketing campaigns, generating **1.2L+ songs across 8 languages** and **80K+ personalized videos (Closeup campaign)**

- Designed pipeline for **lipsync generation, media processing, and FFmpeg-based video stitching**, handling **TB-scale media data**
- Engineered **GPU-based workers + queue-driven autoscaling (AWS ASG)** enabling **200 parallel video jobs**, improving processing throughput by ~3–4x
- Reduced infrastructure cost by ~25–35% through **dynamic scaling and workload-based scheduling**

• Built **RAG-based AI assistants** for enterprise and customer use cases:

- **Amway (internal tool):** Reduced employee knowledge search time by ~40–50% using **FAISS-based semantic retrieval + feedback loops**
- **Himalayan Saffron (customer-facing):** Improved query resolution accuracy by ~30% for product and reference-based queries

• Designed **hybrid LLM architecture (open-source + paid models)** with caching and routing, reducing **inference cost by ~30–40%** and improving **latency by ~20–25%**

• Built **LLM evaluation pipelines** using user feedback (thumbs up/down) and automated checks, reducing hallucination rates and improving response quality by ~25%

• Developed **multimodal AI pipelines** converting **audio, video, and text into structured insights**, enabling downstream analytics and decision-making

• Built and deployed **scalable FastAPI services on AWS (EC2, S3, RDS)** handling high-throughput AI workloads with low-latency response times

• Implemented **monitoring and logging systems** tracking LLM performance, failures, and user feedback, improving system reliability and observability

• Designed **agentic AI workflows** enabling multi-step reasoning, context retention, and automation across complex business processes

PROJECTS

Appointment-Agent | LLM-based Scheduling System [github]

• Built a **no-code, API-first appointment assistant** using LLMs, enabling automated booking, availability checks, and natural language scheduling (e.g., “tomorrow evening”)

• Designed **agent-based architecture (LangChain)** with tool calling and conversation memory for multi-step reasoning and stateful interactions

• Integrated **RAG (PDF knowledge base)** to handle contextual queries (e.g., clinic info, doctor details) with improved response relevance

• Developed **scalable FastAPI services** with validation workflows, ensuring accurate scheduling (name/date/time) and preventing hallucinations

SKILLS

AI / ML: GenAI, LLMs, RAG, NLP, Agentic AI, Semantic Search

LLM & Orchestration: LangChain, Prompt Engineering, LLM Evaluation, RAG Optimization, Agentic Workflows

Backend & MLOps: FastAPI, Docker, CI/CD (GitHub Actions), API Serving, Monitoring & Logging

Cloud & Data: AWS (EC2, S3, RDS, ECR), RunPod (GPU deployments) , MySQL, MongoDB, FAISS, Pinecone

LLM Providers: Groq, OpenAI, Google (Vertex AI), Segmind, ElevenLabs

Model Deployment: Deployed **14B parameter models** on **NVIDIA A100 & H100 GPUs**, optimizing inference performance and throughput

EDUCATION

University of Delhi — New Delhi, India

B.Sc. (Hons.) Mathematics with Computer Science | CGPA: 7.73

Jul 2020 – Jun 2023

Jawahar Navodaya Vidyalaya (JNV), Bageshwar

Class VI – XII | CGPA: 9.5

Jul 2013 – Jun 2020